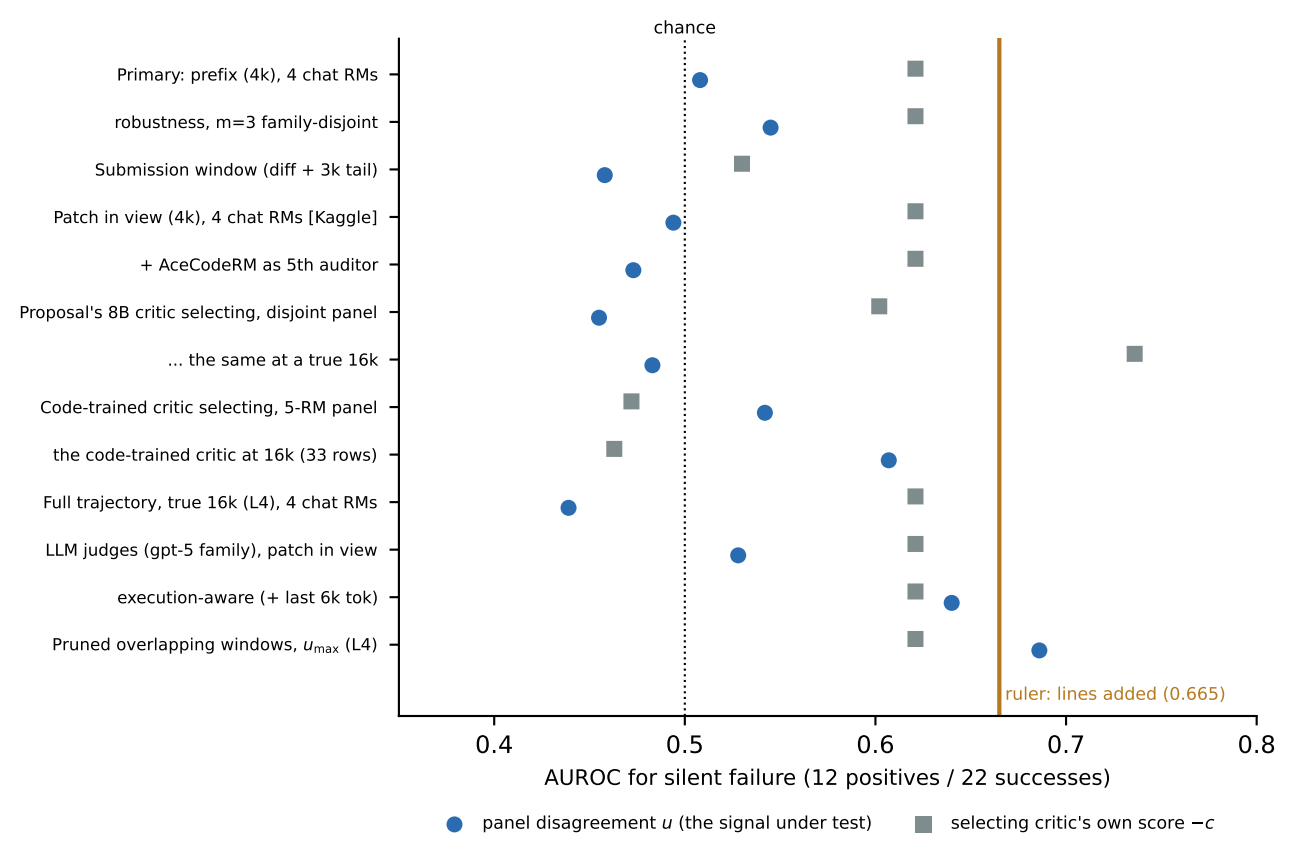
chance
Primary: prefix (4k), 4 chat RMs
robustness, m=3 family-disjoint
Submission window (diff + 3k tail)
Patch in view (4k), 4 chat RMs [Kaggle]
+ AceCodeRM as 5th auditor
Proposal's 8B critic selecting, disjoint panel
... the same at a true 16k
Code-trained critic selecting, 5-RM panel
the code-trained critic at 16k (33 rows)
Full trajectory, true 16k (L4), 4 chat RMs
LLM judges (gpt-5 family), patch in view
execution-aware (+ last 6k tok)
Pruned overlapping windows, $u_{\max}$ (L4)
ruler: lines added (0.665)
0.4
0.5
0.6
0.7
0.8
AUROC for silent failure (12 positives / 22 successes)
panel disagreement u (the signal under test)
selecting critic's own score $-c$