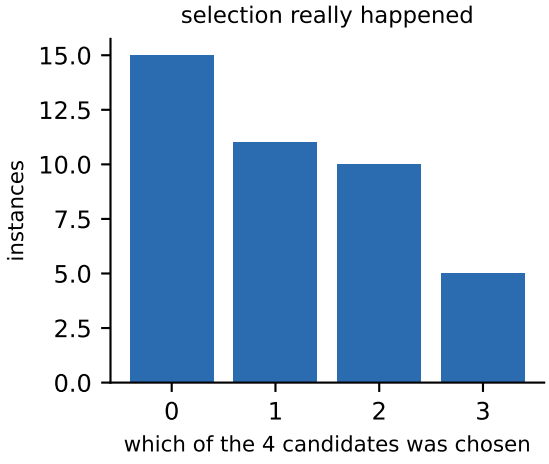
selection really happened
instances
15.0
12.5
10.0
7.5
5.0
2.5
0.0
0
1
2
3
which of the 4 candidates was chosen

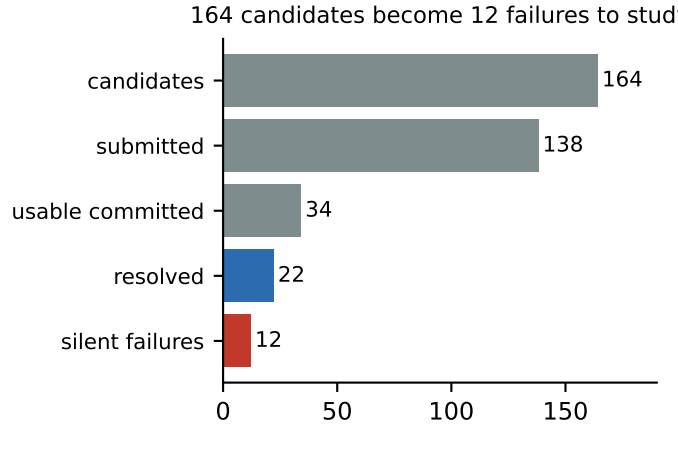
164 candidates become 12 failures to study
candidates
submitted
usable committed
resolved
silent failures
164
138
34
22
12
0
50
100
150

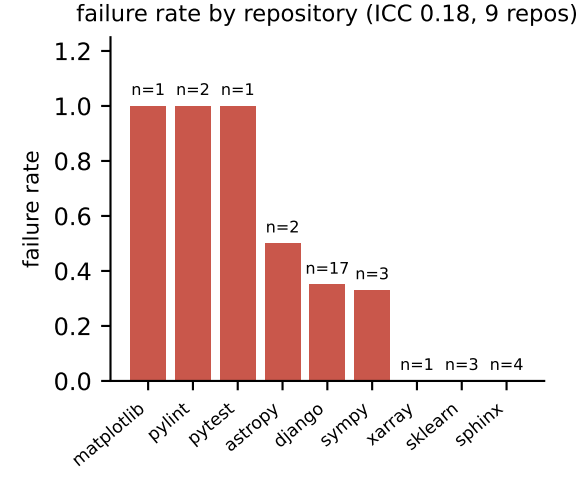
failure rate by repository (ICC 0.18, 9 repos)
failure rate
1.2
1.0
0.8
0.6
0.4
0.2
0.0
n=1
n=2
n=1
n=2
n=17
n=3
n=1
n=3
n=4
matplotlib
pylint
pytest
astropy
django
sympy
xarray
sklearn
sphinx

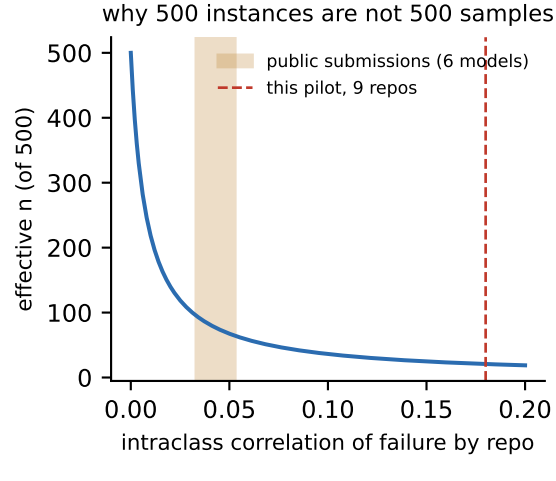
why 500 instances are not 500 samples
effective n (of 500)
500
400
300
200
100
0
public submissions (6 models)
this pilot, 9 repos
0.00
0.05
0.10
0.15
0.20
intraclass correlation of failure by repo