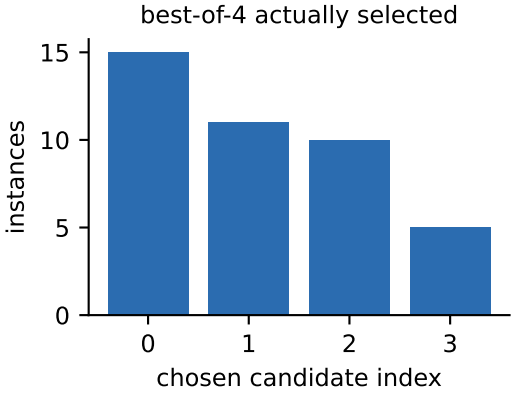
best-of-4 actually selected
instances
15
10
5
0
0
1
2
3
chosen candidate index

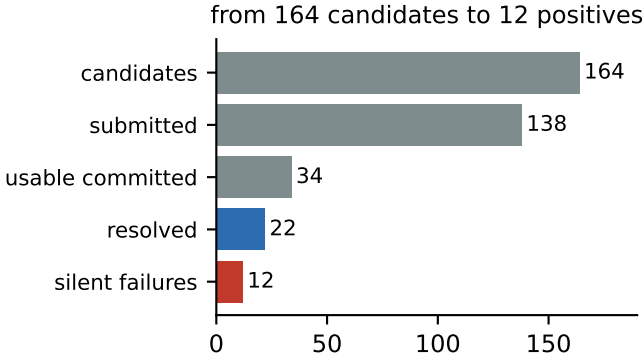
from 164 candidates to 12 positives
candidates
submitted
usable committed
resolved
silent failures
164
138
34
22
12
0
50
100
150

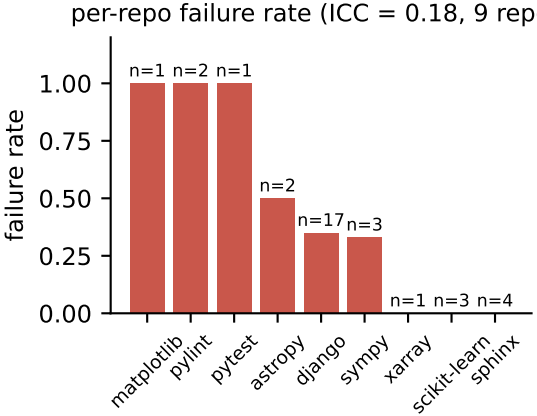
per-repo failure rate (ICC = 0.18, 9 repos)
failure rate
1.00
0.75
0.50
0.25
0.00
n=1
n=2
n=1
n=2
n=17
n=3
n=1
n=3
n=4
matplotlib
pylint
pytest
astropy
django
sympy
xarray
scikit-learn
sphinx

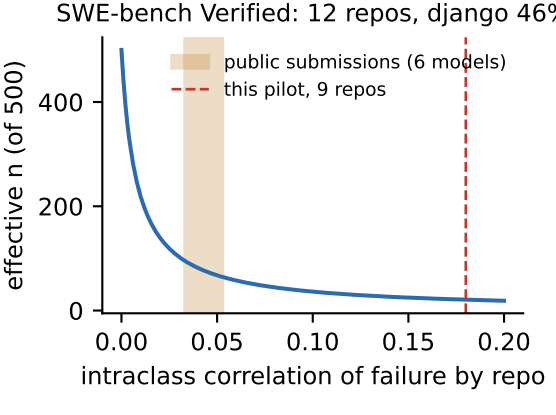
SWE-bench Verified: 12 repos, django 46%
public submissions (6 models)
this pilot, 9 repos
effective n (of 500)
400
200
0
0.00
0.05
0.10
0.15
0.20
intraclass correlation of failure by repo