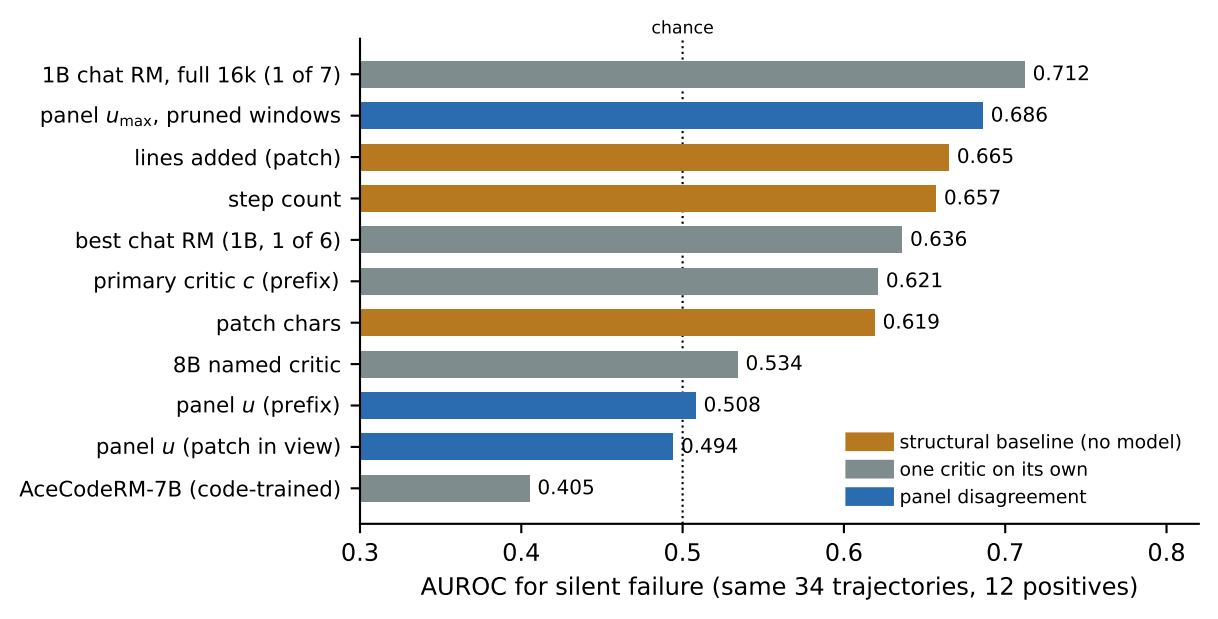
chance
1B chat RM, full 16k (1 of 7)
panel $u_{max}$, pruned windows
lines added (patch)
step count
best chat RM (1B, 1 of 6)
primary critic $c$ (prefix)
patch chars
8B named critic
panel $u$ (prefix)
panel $u$ (patch in view)
AceCodeRM-7B (code-trained)
0.712
0.686
0.665
0.657
0.636
0.621
0.619
0.534
0.508
0.494
0.405
structural baseline (no model)
one critic on its own
panel disagreement
0.3
0.4
0.5
0.6
0.7
0.8
AUROC for silent failure (same 34 trajectories, 12 positives)